

## Краткий вводный обзор Python-библиотек для data science

Сафин Руслан, технический директор Vyndyusoft™



@razonrus

### Для кого этот доклад?

## Анализ данных

### Pandas

```
import pandas as pd
```

### IBM HR Analytics Employee Attrition & Performance

Predict attrition of your valuable employees

<https://www.kaggle.com/pavansubhasht/ibm-hr-analytics-attrition-dataset>

```
df = pd.read_csv('data/WA_Fn-UseC_-HR-Employee-Attrition.csv')
```

```
df.head()
```

|   | Age | Attrition | BusinessTravel    | DailyRate | Department             | DistanceFromHome | Education | EducationField | EmployeeCount | EmployeeNurr |
|---|-----|-----------|-------------------|-----------|------------------------|------------------|-----------|----------------|---------------|--------------|
| 0 | 41  | Yes       | Travel_Rarely     | 1102      | Sales                  | 1                | 2         | Life Sciences  | 1             | 1            |
| 1 | 49  | No        | Travel_Frequently | 279       | Research & Development | 8                | 1         | Life Sciences  | 1             | 2            |
| 2 | 37  | Yes       | Travel_Rarely     | 1373      | Research & Development | 2                | 2         | Other          | 1             | 4            |

|   | Age | Attrition | BusinessTravel    | DailyRate | Department             | DistanceFromHome | Education | EducationField | EmployeeCount | EmployeeNurr |
|---|-----|-----------|-------------------|-----------|------------------------|------------------|-----------|----------------|---------------|--------------|
| 3 | 33  | No        | Travel_Frequently | 1392      | Research & Development | 3                | 4         | Life Sciences  | 1             | 5            |
| 4 | 27  | No        | Travel_Rarely     | 591       | Research & Development | 2                | 1         | Medical        | 1             | 7            |

5 rows × 35 columns

```
df.columns
```

```
Index([u'Age', u'Attrition', u'BusinessTravel', u'DailyRate', u'Department',
      u'DistanceFromHome', u'Education', u'EducationField', u'EmployeeCount',
      u'EmployeeNumber', u'EnvironmentSatisfaction', u'Gender', u'HourlyRate',
      u'JobInvolvement', u'JobLevel', u'JobRole', u'JobSatisfaction',
      u'MaritalStatus', u'MonthlyIncome', u'MonthlyRate',
      u'NumCompaniesWorked', u'Over18', u'OverTime', u'PercentSalaryHike',
      u'PerformanceRating', u'RelationshipSatisfaction', u'StandardHours',
      u'StockOptionLevel', u'TotalWorkingYears', u'TrainingTimesLastYear',
      u'WorkLifeBalance', u'YearsAtCompany', u'YearsInCurrentRole',
      u'YearsSinceLastPromotion', u'YearsWithCurrManager'],
      dtype='object')
```

```
df.shape
```

```
(1470, 35)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
Age                1470 non-null int64
Attrition          1470 non-null object
BusinessTravel     1470 non-null object
DailyRate          1470 non-null int64
Department         1470 non-null object
DistanceFromHome   1470 non-null int64
Education          1470 non-null int64
EducationField     1470 non-null object
EmployeeCount      1470 non-null int64
EmployeeNumber     1470 non-null int64
EnvironmentSatisfaction 1470 non-null int64
Gender             1470 non-null object
HourlyRate         1470 non-null int64
JobInvolvement     1470 non-null int64
JobLevel           1470 non-null int64
JobRole            1470 non-null object
JobSatisfaction    1470 non-null int64
MaritalStatus      1470 non-null object
MonthlyIncome      1470 non-null int64
MonthlyRate        1470 non-null int64
NumCompaniesWorked 1470 non-null int64
Over18             1470 non-null object
OverTime           1470 non-null object
PercentSalaryHike  1470 non-null int64
PerformanceRating  1470 non-null int64
RelationshipSatisfaction 1470 non-null int64
StandardHours      1470 non-null int64
StockOptionLevel   1470 non-null int64
TotalWorkingYears  1470 non-null int64
TrainingTimesLastYear 1470 non-null int64
WorkLifeBalance    1470 non-null int64
YearsAtCompany     1470 non-null int64
```

```

YearsInCurrentRole      1470 non-null int64
YearsSinceLastPromotion  1470 non-null int64
YearsWithCurrManager     1470 non-null int64
dtypes: int64(26), object(9)
memory usage: 402.0+ KB

```

```
df.describe()
```

```

.dataframe thead th {
    text-align: left;
}

.dataframe tbody tr th {
    vertical-align: top;
}

```

|              | Age         | DailyRate   | DistanceFromHome | Education   | EmployeeCount | EmployeeNumber | EnvironmentSatisfaction | HourlyRate  |
|--------------|-------------|-------------|------------------|-------------|---------------|----------------|-------------------------|-------------|
| <b>count</b> | 1470.000000 | 1470.000000 | 1470.000000      | 1470.000000 | 1470.0        | 1470.000000    | 1470.000000             | 1470.000000 |
| <b>mean</b>  | 36.923810   | 802.485714  | 9.192517         | 2.912925    | 1.0           | 1024.865306    | 2.721769                | 65.891156   |
| <b>std</b>   | 9.135373    | 403.509100  | 8.106864         | 1.024165    | 0.0           | 602.024335     | 1.093082                | 20.329428   |
| <b>min</b>   | 18.000000   | 102.000000  | 1.000000         | 1.000000    | 1.0           | 1.000000       | 1.000000                | 30.000000   |
| <b>25%</b>   | 30.000000   | 465.000000  | 2.000000         | 2.000000    | 1.0           | 491.250000     | 2.000000                | 48.000000   |
| <b>50%</b>   | 36.000000   | 802.000000  | 7.000000         | 3.000000    | 1.0           | 1020.500000    | 3.000000                | 66.000000   |
| <b>75%</b>   | 43.000000   | 1157.000000 | 14.000000        | 4.000000    | 1.0           | 1555.750000    | 4.000000                | 83.750000   |
| <b>max</b>   | 60.000000   | 1499.000000 | 29.000000        | 5.000000    | 1.0           | 2068.000000    | 4.000000                | 100.000000  |

8 rows × 26 columns

```
#s = df.describe().loc['std']
```

```
#s[s == 0]
```

```
df.describe(include=['object', 'bool'])
```

```

.dataframe thead th {
    text-align: left;
}

.dataframe tbody tr th {
    vertical-align: top;
}

```

|               | Attrition | BusinessTravel | Department             | EducationField | Gender | JobRole         | MaritalStatus | Over18 | OverTime |
|---------------|-----------|----------------|------------------------|----------------|--------|-----------------|---------------|--------|----------|
| <b>count</b>  | 1470      | 1470           | 1470                   | 1470           | 1470   | 1470            | 1470          | 1470   | 1470     |
| <b>unique</b> | 2         | 3              | 3                      | 6              | 2      | 9               | 3             | 1      | 2        |
| <b>top</b>    | No        | Travel_Rarely  | Research & Development | Life Sciences  | Male   | Sales Executive | Married       | Y      | No       |
| <b>freq</b>   | 1233      | 1043           | 961                    | 606            | 882    | 326             | 673           | 1470   | 1054     |

```
df = df.drop(['Over18', 'EmployeeCount', 'StandardHours'], axis=1)
```

```
df.shape
```

```
(1470, 32)
```

```
df['Attrition'].value_counts()
```

```
No      1233
Yes      237
Name: Attrition, dtype: int64
```

```
df['JobRole'].value_counts()
```

```
Sales Executive      326
Research Scientist    292
Laboratory Technician 259
Manufacturing Director 145
Healthcare Representative 131
Manager              102
Sales Representative   83
Research Director     80
Human Resources       52
Name: JobRole, dtype: int64
```

```
df['JobRole'].value_counts(normalize=True)
```

```
Sales Executive      0.221769
Research Scientist    0.198639
Laboratory Technician 0.176190
Manufacturing Director 0.098639
Healthcare Representative 0.089116
Manager              0.069388
Sales Representative   0.056463
Research Director     0.054422
Human Resources       0.035374
Name: JobRole, dtype: float64
```

```
df[df['Attrition'] == 'Yes'].mean()
```

```
Age                33.607595
DailyRate          750.362869
DistanceFromHome    10.632911
Education           2.839662
EmployeeNumber     1010.345992
EnvironmentSatisfaction 2.464135
HourlyRate         65.573840
JobInvolvement      2.518987
JobLevel           1.637131
JobSatisfaction     2.468354
MonthlyIncome      4787.092827
MonthlyRate        14559.308017
NumCompaniesWorked  2.940928
PercentSalaryHike   15.097046
PerformanceRating   3.156118
RelationshipSatisfaction 2.599156
StockOptionLevel    0.527426
TotalWorkingYears   8.244726
TrainingTimesLastYear 2.624473
WorkLifeBalance     2.658228
YearsAtCompany      5.130802
YearsInCurrentRole  2.902954
YearsSinceLastPromotion 1.945148
```

```
YearsWithCurrManager      2.852321
dtype: float64
```

```
df[df['Attrition'] == 'Yes']['YearsWithCurrManager'].mean()
```

```
2.852320675105485
```

```
df[df['Attrition'] == 'No']['YearsWithCurrManager'].mean()
```

```
4.367396593673966
```

```
df[(df['MaritalStatus'] != 'Married') & (df['BusinessTravel'] == 'Travel_Frequently']]['Attrition'].value_counts(normalize=True)
```

```
No      0.685535
Yes      0.314465
Name: Attrition, dtype: float64
```

```
df['Attrition'].value_counts(normalize=True)
```

```
No      0.838776
Yes      0.161224
Name: Attrition, dtype: float64
```

## Группировка

```
df.groupby(['Attrition'])[['DistanceFromHome', 'YearsWithCurrManager']].describe(percentiles=[])
```

```
.dataframe thead th {
  text-align: left;
}

.dataframe tbody tr th {
  vertical-align: top;
}
```

|           | DistanceFromHome |           |          |     |     |      | YearsWithCurrManager |          |          |     |     |      |
|-----------|------------------|-----------|----------|-----|-----|------|----------------------|----------|----------|-----|-----|------|
|           | count            | mean      | std      | min | 50% | max  | count                | mean     | std      | min | 50% | max  |
| Attrition |                  |           |          |     |     |      |                      |          |          |     |     |      |
| No        | 1233.0           | 8.915653  | 8.012633 | 1.0 | 7.0 | 29.0 | 1233.0               | 4.367397 | 3.594116 | 0.0 | 3.0 | 17.0 |
| Yes       | 237.0            | 10.632911 | 8.452525 | 1.0 | 9.0 | 29.0 | 237.0                | 2.852321 | 3.143349 | 0.0 | 2.0 | 14.0 |

## Сводные таблицы

```
pd.crosstab(df['Attrition'], df['MaritalStatus'])
```

```
.dataframe thead th {
  text-align: left;
}
```

```
.dataframe tbody tr th {
    vertical-align: top;
}
```

| MaritalStatus | Divorced | Married | Single |
|---------------|----------|---------|--------|
| Attrition     |          |         |        |
| No            | 294      | 589     | 350    |
| Yes           | 33       | 84      | 120    |

```
pd.crosstab(df['Attrition'], df['MaritalStatus'], normalize=True)
```

```
.dataframe thead th {
    text-align: left;
}

.dataframe tbody tr th {
    vertical-align: top;
}
```

| MaritalStatus | Divorced | Married  | Single   |
|---------------|----------|----------|----------|
| Attrition     |          |          |          |
| No            | 0.200000 | 0.400680 | 0.238095 |
| Yes           | 0.022449 | 0.057143 | 0.081633 |

```
df.pivot_table(['DailyRate', 'Education', 'TotalWorkingYears'],
                ['Department'], aggfunc='mean')
```

```
.dataframe thead th {
    text-align: left;
}

.dataframe tbody tr th {
    vertical-align: top;
}
```

|                        | DailyRate  | Education | TotalWorkingYears |
|------------------------|------------|-----------|-------------------|
| Department             |            |           |                   |
| Human Resources        | 751.539683 | 2.968254  | 11.555556         |
| Research & Development | 806.851197 | 2.899063  | 11.342352         |
| Sales                  | 800.275785 | 2.934978  | 11.105381         |

```
pd.crosstab(df['Attrition'], df['BusinessTravel'], margins=True)
```

```
.dataframe thead th {
    text-align: left;
}

.dataframe tbody tr th {
    vertical-align: top;
}
```

| BusinessTravel | Non-Travel | Travel_Frequently | Travel_Rarely | All  |
|----------------|------------|-------------------|---------------|------|
| Attrition      |            |                   |               |      |
| No             | 138        | 208               | 887           | 1233 |

| BusinessTravel | Non-Travel | Travel_Frequently | Travel_Rarely | All  |
|----------------|------------|-------------------|---------------|------|
| Attrition      |            |                   |               |      |
| Yes            | 12         | 69                | 156           | 237  |
| All            | 150        | 277               | 1043          | 1470 |

## Визуализация данных

### Seaborn и Matplotlib

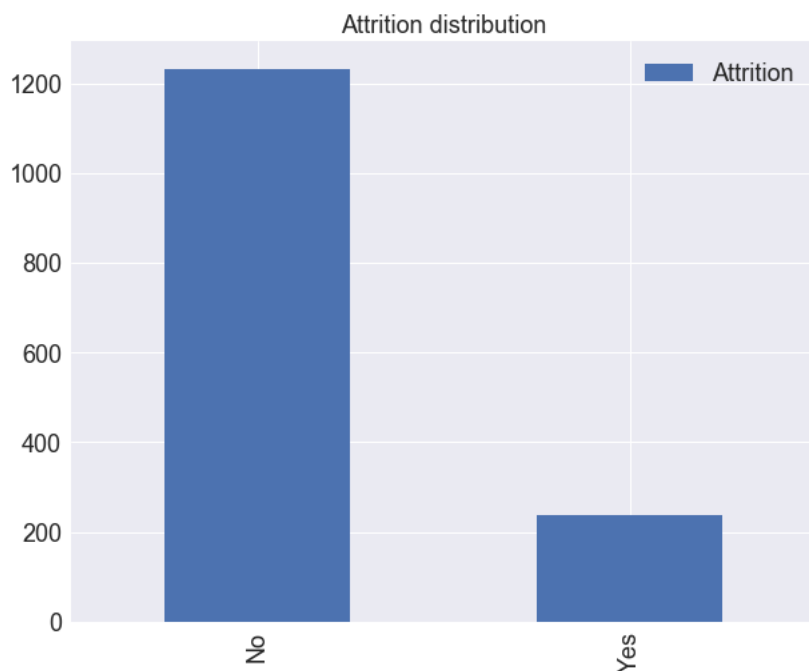
```
%matplotlib inline
import seaborn as sns
import matplotlib.pyplot as plt
#графики в svg выглядят более четкими
%config InlineBackend.figure_format = 'svg'

#увеличим дефолтный размер графиков
from pylab import rcParams
rcParams['figure.figsize'] = 10, 8;
```

```
BIGGER_SIZE = 18

plt.rc('font', size=BIGGER_SIZE)          # controls default text sizes
plt.rc('axes', titlesize=BIGGER_SIZE)     # fontsize of the axes title
plt.rc('axes', labelsizе=BIGGER_SIZE)    # fontsize of the x and y labels
plt.rc('xtick', labelsizе=BIGGER_SIZE)    # fontsize of the tick labels
plt.rc('ytick', labelsizе=BIGGER_SIZE)    # fontsize of the tick labels
plt.rc('legend', fontsize=BIGGER_SIZE)    # legend fontsize
plt.rc('figure', titlesize=BIGGER_SIZE)   # fontsize of the figure title
```

```
df['Attrition'].value_counts().plot(kind='bar', label='Attrition')
plt.legend()
plt.title('Attrition distribution');
```




```
pd.crosstab(df['Attrition'], df['BusinessTravel'], margins=True)
```

```
.dataframe thead th {
  text-align: left;
}
```


```
.dataframe tbody tr th {
    vertical-align: top;
}
```

| BusinessTravel | Non-Travel | Travel_Frequently | Travel_Rarely | All  |
|----------------|------------|-------------------|---------------|------|
| Attrition      |            |                   |               |      |
| No             | 138        | 208               | 887           | 1233 |
| Yes            | 12         | 69                | 156           | 237  |
| All            | 150        | 277               | 1043          | 1470 |

```
sns.countplot(x='BusinessTravel', hue='Attrition', data=df);
```



```
sns.countplot(x='MaritalStatus', hue='Attrition', data=df);
```



```
pd.crosstab(df['Attrition'], df['StockOptionLevel'], margins=True)
```




```
.dataframe thead th {
    text-align: left;
}

.dataframe tbody tr th {
    vertical-align: top;
}
```

| StockOptionLevel | 0   | 1   | 2   | 3  | All  |
|------------------|-----|-----|-----|----|------|
| Attrition        |     |     |     |    |      |
| No               | 477 | 540 | 146 | 70 | 1233 |
| Yes              | 154 | 56  | 12  | 15 | 237  |
| All              | 631 | 596 | 158 | 85 | 1470 |

```
sns.countplot(x='StockOptionLevel', hue='Attrition', data=df);
```



```
#df['Risk'] = ((df['MaritalStatus'] != 'Married') & (df['BusinessTravel'] == 'Travel_Frequently') & (df['StockOptionLevel'] == 0)).astype('int')
```


```
#df[['Risk', 'MaritalStatus', 'BusinessTravel', 'StockOptionLevel']].head(6)
```

```
#pd.crosstab(df['Attrition'], df['Risk'])
```

```
#sns.countplot(x='Risk', hue='Attrition', data=df);
```

### Scatter plot matrix


```
cols = ['DailyRate', 'Age', 'TotalWorkingYears']
sns_plot = sns.pairplot(df[cols], size = 4)
sns_plot.savefig('pairplot.png')
```



Гистограмма и KDE ( [kernel density estimation](#) )


```
sns.distplot(df.Age)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x10f61c88>
```



```
sns.jointplot(df.YearsAtCompany, df.MonthlyIncome, size =10)
```

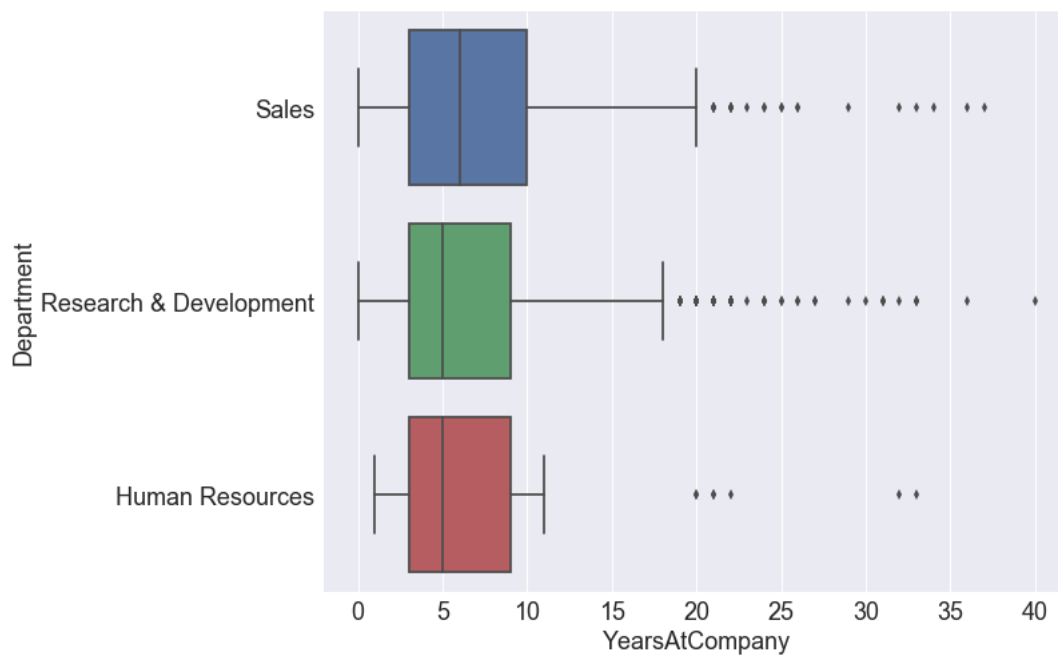
```
<seaborn.axisgrid.JointGrid at 0xfd3f2b0>
```



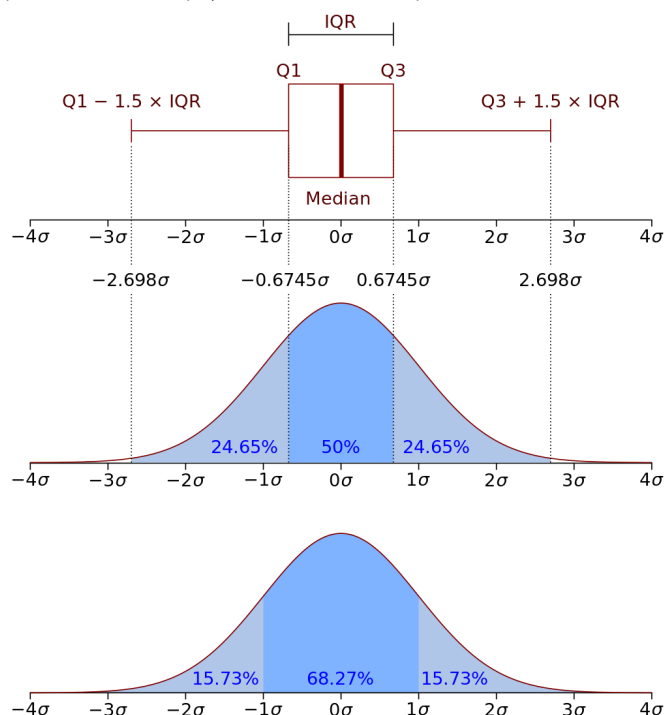
Box plot

```
sns.boxplot(y="Department", x="YearsAtCompany", data=df, orient="h")
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x1117af60>
```



Box plot состоит из коробки (поэтому он и называется box plot), усов и точек. Коробка показывает интерквартильный размах распределения, то есть соответственно 25% (Q1) и 75% (Q3) перцентили. Черта внутри коробки обозначает медиану распределения. С коробкой разобрались, перейдем к усам. Усы отображают весь разброс точек кроме выбросов, то есть минимальные и максимальные значения, которые попадают в промежуток  $(Q1 - 1.5 \times IQR, Q3 + 1.5 \times IQR)$ , где  $IQR = Q3 - Q1$  — интерквартильный размах. Точками на графике обозначаются выбросы (outliers) — те значения, которые не вписываются в



промежуток значений, заданный усами графика.

## Heat map

```
department_ef_mi = df.pivot_table(
    index='Department',
    columns='EducationField',
    values='YearsInCurrentRole',
    aggfunc='mean').fillna(0).applymap(float)
sns.heatmap(department_ef_mi, annot=True, fmt=".1f", linewidths=.5)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x11514b00>
```



```
df['Attrition'] = (df['Attrition']=='Yes').astype('int64')
```

```
corr_matrix = df.corr()
```

```
corr_matrix.head()
```


```
.dataframe thead th {
  text-align: left;
}

.dataframe tbody tr th {
  vertical-align: top;
}
```

|                  | Age       | Attrition | DailyRate | DistanceFromHome | Education | EmployeeNumber | EnvironmentSatisfaction | HourlyRate |
|------------------|-----------|-----------|-----------|------------------|-----------|----------------|-------------------------|------------|
| Age              | 1.000000  | -0.159205 | 0.010661  | -0.001686        | 0.208034  | -0.010145      | 0.010146                | 0.024287   |
| Attrition        | -0.159205 | 1.000000  | -0.056652 | 0.077924         | -0.031373 | -0.010577      | -0.103369               | -0.006846  |
| DailyRate        | 0.010661  | -0.056652 | 1.000000  | -0.004985        | -0.016806 | -0.050990      | 0.018355                | 0.023381   |
| DistanceFromHome | -0.001686 | 0.077924  | -0.004985 | 1.000000         | 0.021042  | 0.032916       | -0.016075               | 0.031131   |
| Education        | 0.208034  | -0.031373 | -0.016806 | 0.021042         | 1.000000  | 0.042070       | -0.027128               | 0.016775   |

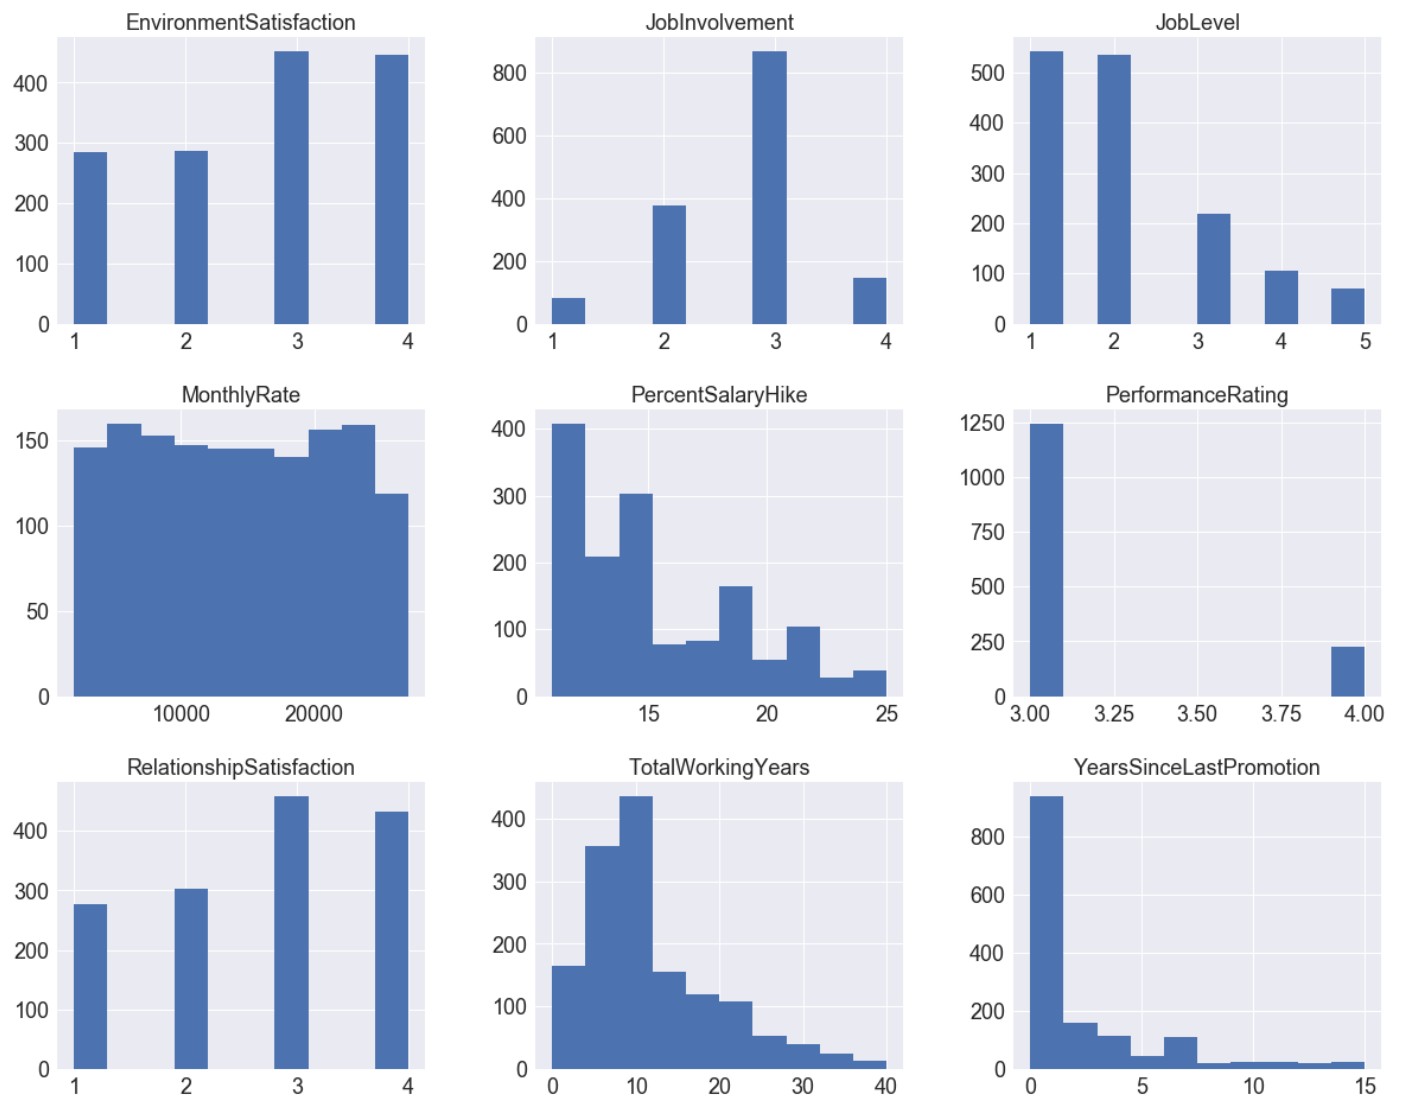
5 rows × 25 columns

```
sns.heatmap(corr_matrix);
```

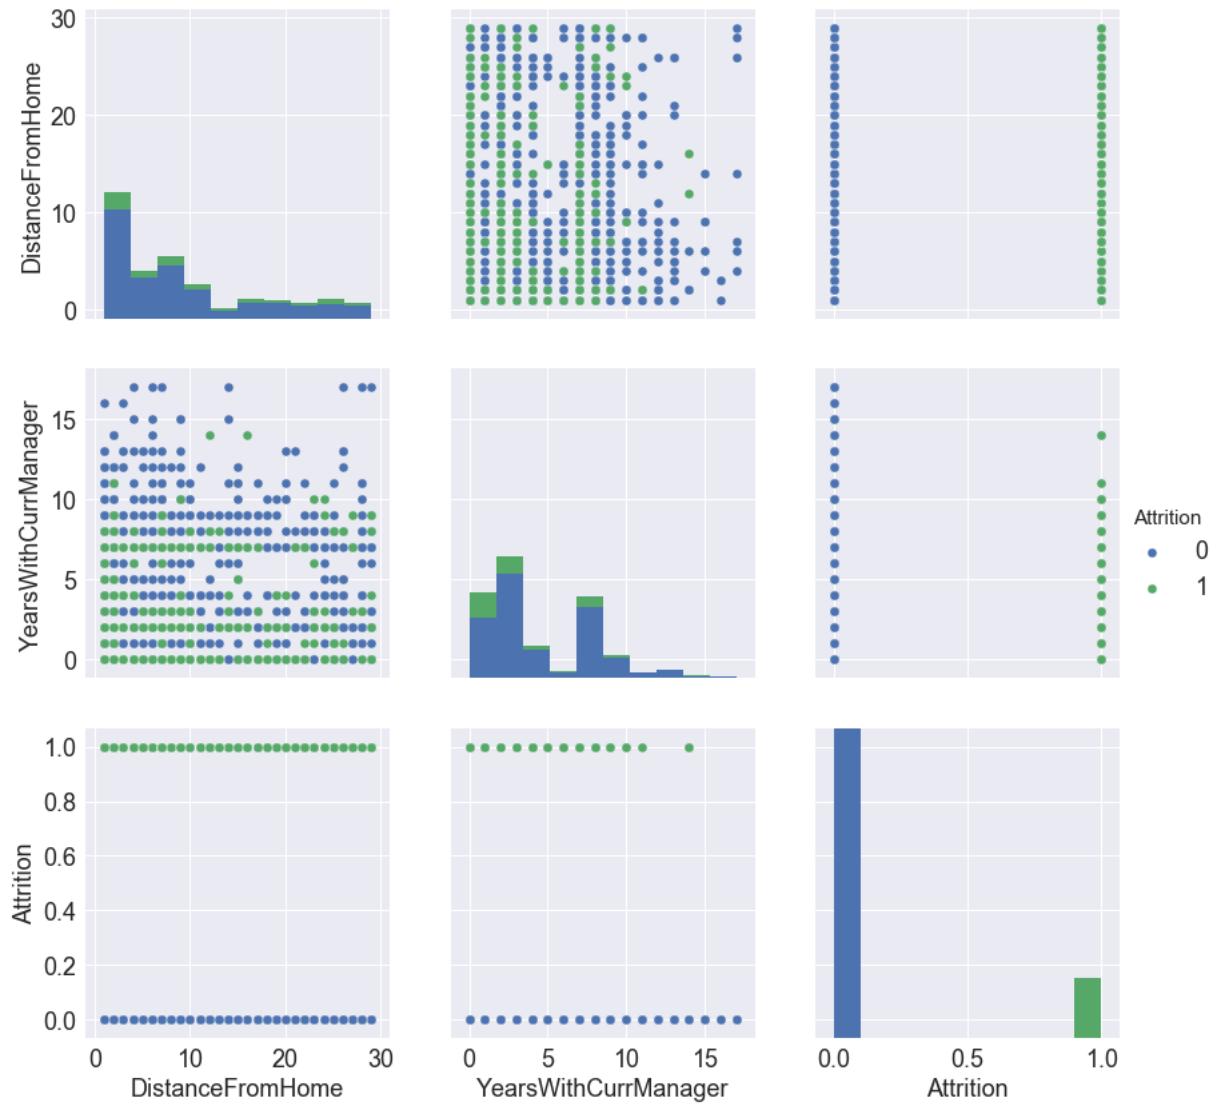


```
features = list(set(df.columns) - set(['BusinessTravel', 'Department', 'EducationField',
                                       'Gender', 'JobRole', 'MaritalStatus',
                                       'Over18', 'OverTime', 'Attrition']))
```

```
df[features[:9]].hist(figsize=(20,16));
```



```
sns.pairplot(df[['DistanceFromHome', 'YearsWithCurrManager', 'Attrition']], hue='Attrition', size = 4);
```



```

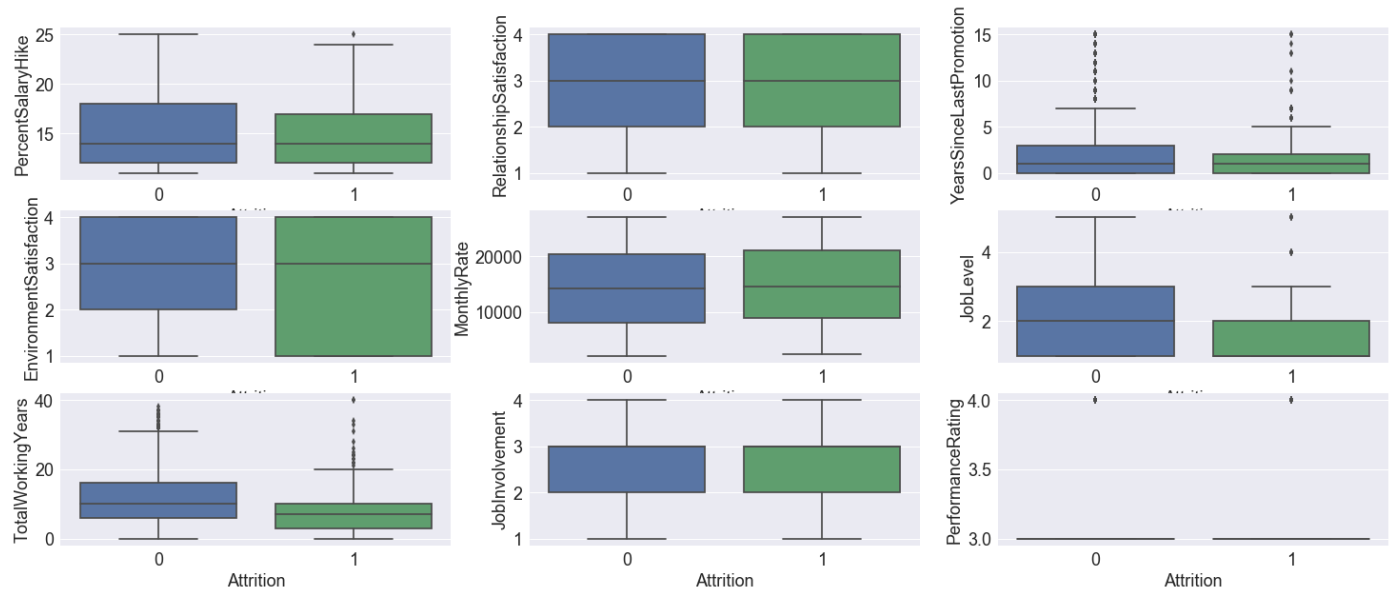
cnt = 3
fig, axes = plt.subplots(nrows=cnt, ncols=cnt, figsize=(25, 10))

for idx, feat in enumerate(features[:9]):
    sns.boxplot(x='Attrition', y=feat, data=df, ax=axes[idx / cnt, idx % cnt])
    axes[idx / cnt, idx % cnt].legend()
    axes[idx / cnt, idx % cnt].set_xlabel('Attrition')
    axes[idx / cnt, idx % cnt].set_ylabel(feat);

```


C:\Anaconda2\lib\site-packages\matplotlib\axes\\_axes.py:545: UserWarning: No labelled objects found. Use label='...' kwarg on individual plots.  
 warnings.warn("No labelled objects found. ")





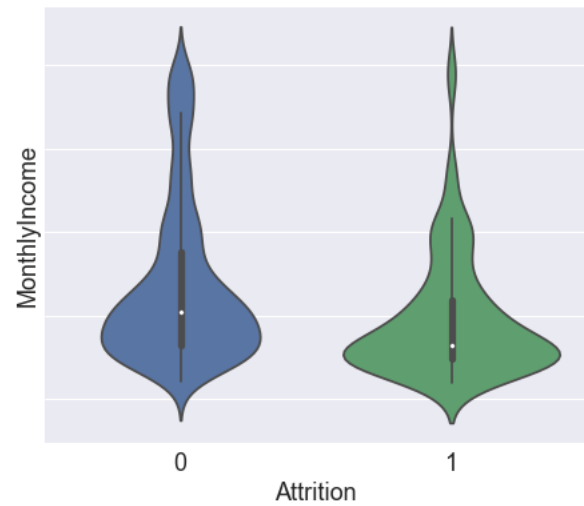

```
_, axes = plt.subplots(1, 2, sharey=True, figsize=(16,6))

sns.boxplot(x='Attrition', y='JobLevel', data=df, ax=axes[0]);
sns.violinplot(x='Attrition', y='JobLevel', data=df, ax=axes[1]);
```



```
_, axes = plt.subplots(1, 2, sharey=True, figsize=(16,6))

sns.boxplot(x='Attrition', y='MonthlyIncome', data=df, ax=axes[0]);
sns.violinplot(x='Attrition', y='MonthlyIncome', data=df, ax=axes[1]);
```



```
#sns.countplot(x='JobLevel', hue='Attrition', data=df);
```

```
#_, axes = plt.subplots(1, 2, sharey=True, figsize=(16,6))

#sns.countplot(x='MaritalStatus', hue='Attrition', data=df, ax=axes[0]);
#sns.countplot(x='Department', hue='Attrition', data=df, ax=axes[1]);
```

### t-SNE (t-distributed Stochastic Neighbor Embedding)

```
from sklearn.manifold import TSNE
from sklearn.preprocessing import StandardScaler
```


```
X = df.drop(['Attrition', 'JobRole', 'BusinessTravel', 'Department', 'EducationField', 'MaritalStatus'], axis=1)
X['Gender'] = pd.factorize(X['Gender'])[0]
X['OverTime'] = pd.factorize(X['OverTime'])[0]

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```


```
%%time
tsne = TSNE(random_state=17)
tsne_representation = tsne.fit_transform(X_scaled) #1min
```

Wall time: 17.4 s

```
plt.scatter(tsne_representation[:, 0], tsne_representation[:, 1]);
```




```
plt.scatter(tsne_representation[:, 0], tsne_representation[:, 1],
            c=df['Attrition'].map({0: 'blue', 1: 'orange'}));
```




```
_, axes = plt.subplots(1, 2, sharey=True, figsize=(16,6))

axes[0].scatter(tsne_representation[:, 0], tsne_representation[:, 1],
                c=df['MaritalStatus'].map({'Married': 'blue', 'Single': 'orange', 'Divorced': 'green'}));
axes[1].scatter(tsne_representation[:, 0], tsne_representation[:, 1],
                c=df['BusinessTravel'].map({'Travel_Rarely': 'blue', 'Travel_Frequently': 'orange', 'Non-Travel': 'green'}));
axes[0].set_title('MaritalStatus');
axes[1].set_title('BusinessTravel');
```



```
_, axes = plt.subplots(1, 2, sharey=True, figsize=(16,6))

axes[0].scatter(tsne_representation[:, 0], tsne_representation[:, 1],
               c=df['StockOptionLevel']);
axes[1].scatter(tsne_representation[:, 0], tsne_representation[:, 1],
               c=df['Education']);
axes[0].set_title('StockOptionLevel');
axes[1].set_title('Education');
```



## Машинное обучение

Scikit-learn. Деревья решений и метод ближайших соседей

```
df['Department'] = pd.factorize(df['Department'])[0]
df['Gender'] = pd.factorize(df['Gender'])[0]
df['JobRole'] = pd.factorize(df['JobRole'])[0]
df['MaritalStatus'] = pd.factorize(df['MaritalStatus'])[0]
df['OverTime'] = pd.factorize(df['OverTime'])[0]
df['EducationField'] = pd.factorize(df['EducationField'])[0]
df['BusinessTravel'] = pd.factorize(df['BusinessTravel'])[0]
```

```
y = df['Attrition']
```

```
df.drop(['Attrition'], axis=1, inplace=True)
```

```
y.value_counts(normalize=True)
```

```
0    0.838776
1    0.161224
Name: Attrition, dtype: float64
```

```
from sklearn.model_selection import train_test_split, StratifiedKFold
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
```

```
X_train, X_holdout, y_train, y_holdout = train_test_split(df.values, y, test_size=0.3,
random_state=17)

tree = DecisionTreeClassifier(max_depth=5, random_state=17)
knn = KNeighborsClassifier(n_neighbors=10)

tree.fit(X_train, y_train)
knn.fit(X_train, y_train)
```

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=1, n_neighbors=10, p=2,
                    weights='uniform')
```

```
from sklearn.metrics import accuracy_score
```

```
tree_pred = tree.predict(X_holdout)
accuracy_score(y_holdout, tree_pred)
```

```
0.83673469387755106
```

```
knn_pred = knn.predict(X_holdout)
accuracy_score(y_holdout, knn_pred)
```

```
0.8344671201814059
```

```
from sklearn.model_selection import GridSearchCV, cross_val_score
```

```
tree_params = {'max_depth': range(1,4), 'max_features': range(10,20)}
```

```
tree_grid = GridSearchCV(tree, tree_params, cv=5, n_jobs=-1, verbose=True)
```

```
%%time
tree_grid.fit(X_train, y_train) #12sec
```

```
Fitting 5 folds for each of 30 candidates, totalling 150 fits
Wall time: 5.06 s
```

```
[Parallel(n_jobs=-1)]: Done 150 out of 150 | elapsed: 4.6s finished
```

```
GridSearchCV(cv=5, error_score='raise',
             estimator=DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=5,
                                              max_features=None, max_leaf_nodes=None,
                                              min_impurity_split=1e-07, min_samples_leaf=1,
                                              min_samples_split=2, min_weight_fraction_leaf=0.0,
                                              presort=False, random_state=17, splitter='best'),
             fit_params={}, iid=True, n_jobs=-1,
             param_grid={'max_features': [10, 11, 12, 13, 14, 15, 16, 17, 18, 19], 'max_depth': [1, 2, 3]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
             scoring=None, verbose=True)
```

```
tree_grid.best_params_
```

```
{'max_depth': 3, 'max_features': 13}
```

```
tree_grid.best_score_
```

```
0.86297376093294464
```

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
```

```
knn_pipe = Pipeline([('scaler', StandardScaler()), ('knn', KNeighborsClassifier(n_jobs=-1))])
```

```
knn_params = {'knn__n_neighbors': range(1, 10)}
```

```
knn_grid = GridSearchCV(knn_pipe, knn_params,
                        cv=5, n_jobs=-1,
                        verbose=True)
```

```
%time
knn_grid.fit(X_train, y_train) #10sec
```

```
Fitting 5 folds for each of 9 candidates, totalling 45 fits
Wall time: 7.97 s
```

```
[Parallel(n_jobs=-1)]: Done 45 out of 45 | elapsed: 7.6s finished
C:\Anaconda2\lib\site-packages\sklearn\utils\validation.py:429: DataConversionWarning: Data with input dtype int64 was converted to float64 by StandardScaler.
warnings.warn(msg, _DataConversionWarning)
```

```
GridSearchCV(cv=5, error_score='raise',
             estimator=Pipeline(steps=[('scaler', StandardScaler(copy=True, with_mean=True, with_std=True)), ('knn',
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=-1, n_neighbors=5, p=2,
                    weights='uniform'))]),
             fit_params={}, iid=True, n_jobs=-1,
             param_grid={'knn__n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9]},
```

```
pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
scoring=None, verbose=True)
```

```
knn_grid.best_params_, knn_grid.best_score_
```

```
{'knn__n_neighbors': 8}, 0.84839650145772594)
```

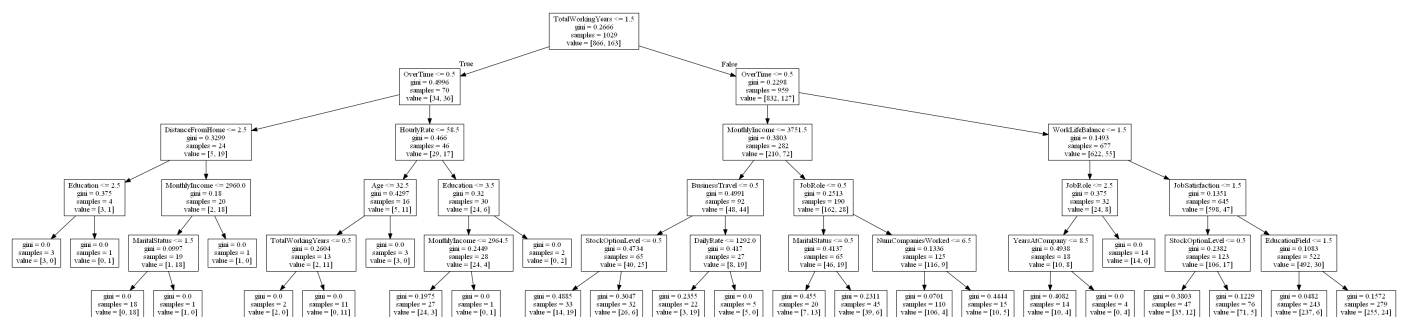
```
import numpy as np
from sklearn.tree import export_graphviz
```

```
#!pip install pydotplus
```

```
#!pip install graphviz
```

```
from sklearn import tree
from IPython.display import Image
import pydotplus
```

```
dot_data = tree.export_graphviz(tree_grid.estimator, feature_names=df.columns, out_file=None)
graph = pydotplus.graph_from_dot_data(dot_data)
graph.write_pdf('df_train.pdf')
graph.write_png('df_train.png')
Image(graph.create_png())
```



## Случайный лес

```
from sklearn.ensemble import RandomForestClassifier
```

```
forest = RandomForestClassifier(n_estimators=100, n_jobs=-1, random_state=17)
print(np.mean(cross_val_score(forest, X_train, y_train, cv=5))) # 0.859
```

```
0.859131853286
```

```
forest_params = {'max_depth': range(10,11),
'max_features': range(10,15)}
```

```
forest_grid = GridSearchCV(forest, forest_params,
cv=5, n_jobs=-1,
verbose=True)
```

```
%%time
forest_grid.fit(X_train, y_train) #50sec
```

Fitting 5 folds for each of 5 candidates, totalling 25 fits

[Parallel(n\_jobs=-1)]: Done 25 out of 25 | elapsed: 13.9s finished

Wall time: 15 s

```
GridSearchCV(cv=5, error_score='raise',
             estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                              max_depth=None, max_features='auto', max_leaf_nodes=None,
                                              min_impurity_split=1e-07, min_samples_leaf=1,
                                              min_samples_split=2, min_weight_fraction_leaf=0.0,
                                              n_estimators=100, n_jobs=-1, oob_score=False, random_state=17,
                                              verbose=0, warm_start=False),
             fit_params={}, iid=True, n_jobs=-1,
             param_grid={'max_features': [10, 11, 12, 13, 14], 'max_depth': [10]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
             scoring=None, verbose=True)
```

```
forest_grid.best_params_, forest_grid.best_score_ #0.864
```

```
({'max_depth': 10, 'max_features': 12}, 0.86394557823129248)
```

## Логистическая регрессия

```
from sklearn.linear_model import LogisticRegression
```

```
%%time
logit = LogisticRegression(n_jobs=-1, random_state=7)
logit.fit(X_train, y_train)
print(round(logit.score(X_train, y_train), 3), round(logit.score(X_holdout, y_holdout), 3))
```

```
(0.869, 0.846)
Wall time: 47 ms
```

```
def visualize_coefficients(classifier, feature_names, n_top_features=25):
    # get coefficients with large absolute values
    coef = classifier.coef_.ravel()
    positive_coefficients = np.argsort(coef)[-n_top_features:]
    negative_coefficients = np.argsort(coef)[:n_top_features]
    interesting_coefficients = np.hstack([negative_coefficients, positive_coefficients])
    # plot them
    plt.figure(figsize=(15, 5))
    colors = ["red" if c < 0 else "blue" for c in coef[interesting_coefficients]]
    plt.bar(np.arange(2 * n_top_features), coef[interesting_coefficients], color=colors)
    feature_names = np.array(feature_names)
    plt.xticks(np.arange(1, 1 + 2 * n_top_features), feature_names[interesting_coefficients], rotation=60, ha="right", size=10);
```

```
#def plot_grid_scores(grid, param_name):
#    plt.plot(grid.param_grid[param_name], grid.cv_results_['mean_train_score'],
```



```
visualize_coefficients(logit, df.columns)
```



<https://habr.com/company/ods/blog/322626/>

Вопросы?